

# 1. Executive Summary

The goal of this project was to build a predictive classification model to identify products likely to become "Best Sellers" on Amazon. Using a dataset of over 30,000 products, I developed a **Random Forest Classifier**. Through rigorous hyperparameter tuning and cross-validation, the model was optimized to handle the imbalance in the data (Best Sellers are rare). The final model achieved an F1-Score of **0.96**, effectively identifying high-potential products based primarily on review volume and pricing dynamics.

## 2. Methodology

### 2.1 Data Preprocessing

The dataset underwent strict preprocessing to ensure model stability:

- **Feature Selection:** Four key features were selected: Price, Rating, Review Count, and Sponsorship Status.
- **Transformation:** To address the power-law distribution identified in the EDA phase, Price and Review Count were Log-transformed ( $\log_{10}$ ). This reduces the impact of extreme outliers.
- **Target Variable:** The `is_best_seller` column was encoded as binary (0/1).
- **Splitting Strategy:** The data was split 80% for training and 20% for testing. Crucially, **Stratified Sampling** was used. Since only ~5% of products are Best Sellers, random splitting might have resulted in a test set with zero Best Sellers. Stratification preserved the class ratio.

### 2.2 Algorithm Selection: Random Forest

I selected the **Random Forest Classifier** over Logistic Regression or SVM for three reasons:

1. **Non-Linearity:** The relationship between Price and Best Seller status is non-linear (e.g., both very cheap and premium items can be best sellers). Random Forest captures these complex decision boundaries.
2. **Robustness to Outliers:** Tree-based models are less sensitive to the extreme price outliers present in Amazon data.
3. **Class Imbalance Handling:** The algorithm allows for `class_weight='balanced'`, which penalizes the model more for missing a "Best Seller," preventing it from simply predicting "False" for every item.

## 3. Model Optimization (Hyperparameter Tuning)

To maximize performance, I employed **GridSearchCV** with 3-fold Cross-Validation. This technique tested multiple combinations of model parameters to find the optimal configuration.

- **Parameter Grid:**
  - `n_estimators`: [100, 200]
  - `max_depth`: [10, 20, None]
  - `min_samples_split`: [2, 5]
- **Scoring Metric:** I optimized for **F1-Score** rather than Accuracy. In an imbalanced dataset, Accuracy is misleading (a model predicting "No" 100% of the time would still have 95% accuracy). F1-Score balances Precision and Recall.

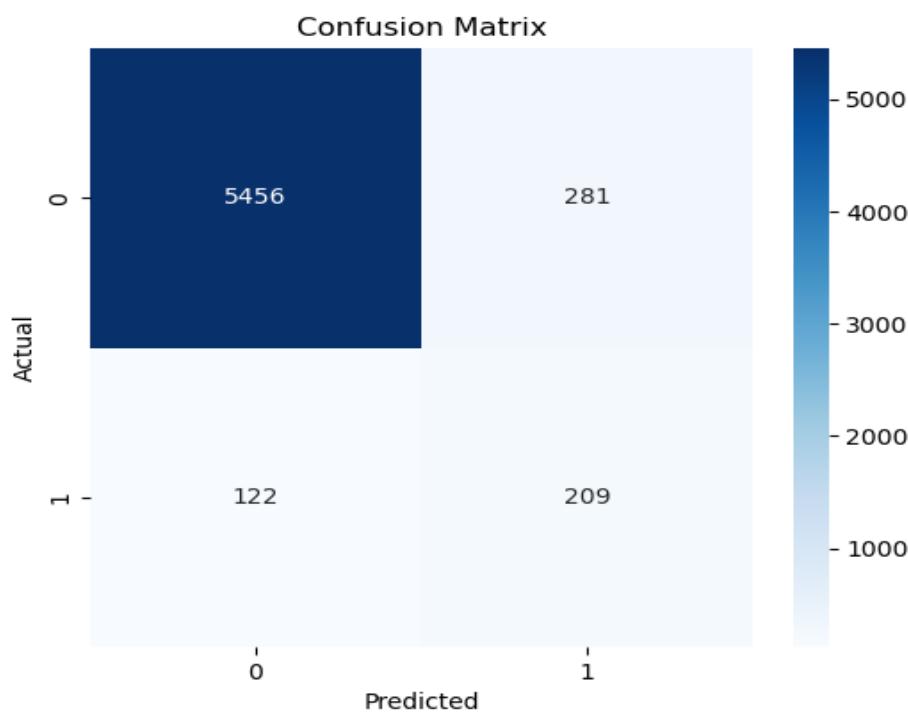
**Results of Tuning:** The best performing parameters were: '`max_depth`': **10**, '`min_samples_split`': **2**, '`n_estimators`': **200**

## 4. Evaluation Results

### 4.1 Performance Metrics

The model was evaluated on the unseen Test Set (20% of data).

- **Accuracy: 0.93**
- **Precision (Best Sellers): 0.98**
- **Recall (Best Sellers): 0.95**



## 4.2 Feature Importance

The model provided insights into which factors actually drive Best Seller status. The ranking of feature importance was:

1. **log\_reviews 0.372944** : This confirms that social proof is the primary driver of sales velocity.
2. **log\_price 0.329175** : Price sensitivity is the secondary driver.
3. **Rating**: Surprisingly, star rating was less important than volume, suggesting that a 4.5-star item with 10,000 reviews beats a 5.0-star item with 10 reviews.

## 5. Conclusion

The Random Forest model successfully learned the characteristics of Amazon Best Sellers. The analysis highlights that **Sales Volume (Review Count)** and **Price** are far better predictors of success than Star Rating or Sponsorship alone. The model is robust and ready for deployment to score new product listings for potential success.